Lecture 5 : Pattern Recognition(1)

Special Topics in Image Processing

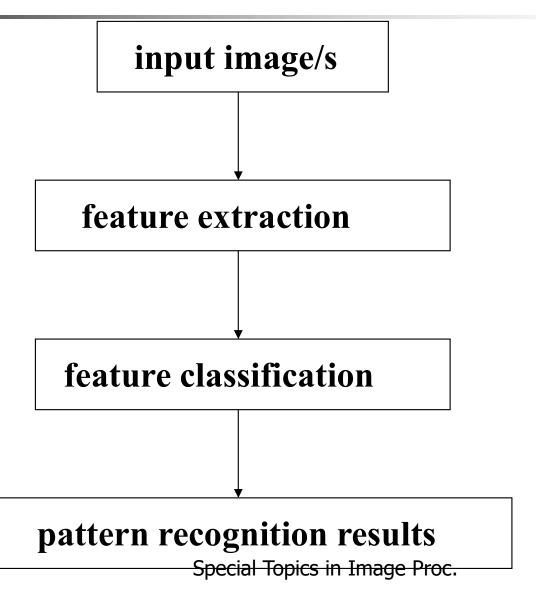


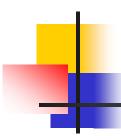
Contents

- Reference:
 - Shapiro Chap 4
 - Duda & Hart (2nd edition) Chap 2,3,10
- Features
- Principle Components Analysis (PCA)
- Bayes Decision Theory
- Gaussian Normal Density



Pattern Recognition Basic Steps



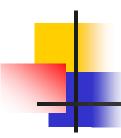


Feature Extraction

Linear features
 eg. Lines, curves

2D features
 eg. Texture

Feature of features



Feature Extraction

- A good feature should be invariant to changes in
 - 1. scale
 - 2. rotation
 - 3. translation
 - 4. illumination



1. Spatial Feature Extraction

- Feature Vectors & Feature Spaces
- Amplitude Features
- Binary Object Features
- Histogram Features



1.1 Feature Vectors and Feature Space

Feature vectors

- one method to represent an image, or part of an image, by finding measurements on a set of features.
- An n-dimensional vector that contain these measurements.

Feature spaces

- a mathematical abstraction associated with the feature vector.
- Also n-dimensional.

Compare two feature vectors

difference or similarity

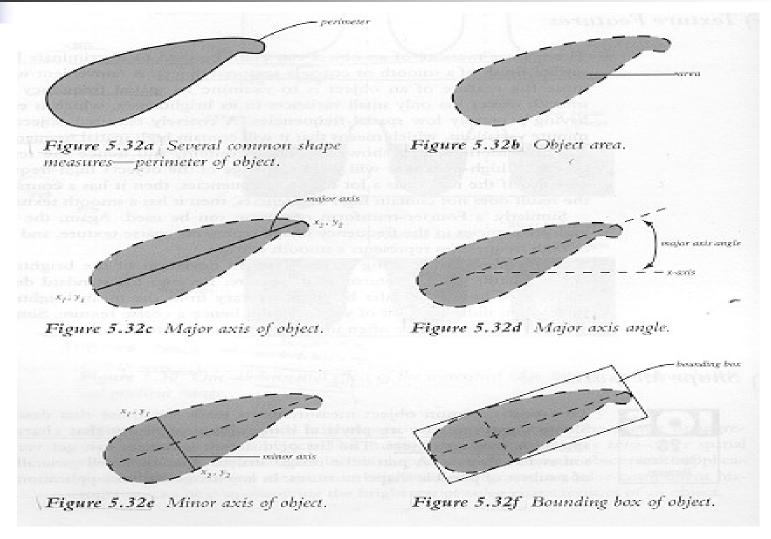
1.2 Amplitude Features

- The simplest and perhaps the most useful.
- Reflectivity, Transmissitivity, color, Multispectral response.
- Can be extracted easily by intensity window slicing or by the more general point transformation
- **E**x)
 - X-ray image represent the absorption characteristics of the body masses.
 - Infrared image
 represents temperature, which facilitates the segmentation of clouds
 from terrain.
 - Radar Image
 represents the radar cross section, which determine the size of the
 object being imaged.

1.3 Binary Object Features

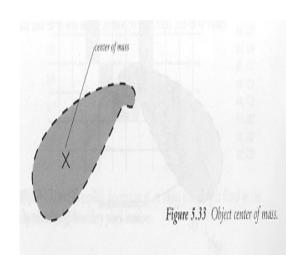
- Area
 - indicates the relative size of the object
- Center of area
 - row and column coordinates of the center of area for the object
- Axis of least second moment
 - provides information about the object's orientation
- Perimeter
 - help us to locate it in space
 - provide us with information about the shape of the object
- Thinness
 - be used as a measure of roundness.
- Euler number
 - tellsous how many closed relieves the relief contains.

Feature Example (1)

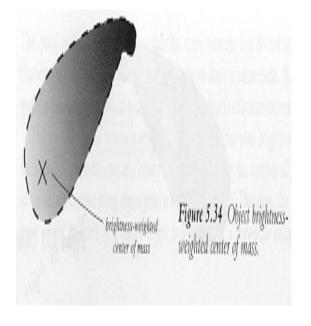




- Center of mass
- Center of Massx = Sum of object's x-pixel coordinates/ Number of pixels in object



Gray Center(brightness center of mass)





1.4 Histogram Features[1]

$$p_u(x) \equiv \text{Prob}[u = x]$$

$$\approx \frac{\text{No. of pixels with gray level } x}{\text{Total no. of pixels in the region}}$$

- A plot of the gray-level values versus the number of pixels at that value
- The shape of the histogram provides us with information about the nature of the image

Moments :
$$m_{i} = E[u^{i}] = \sum_{x=0}^{L-1} x^{i} p_{u}(x)$$
Absolute moments : $\hat{m}_{i} = E[|u|^{i}] = \sum_{x=0}^{L-1} |x|^{i} p_{u}(x)$

Central moments : $\mu_{i} = E\{[\mu - E(u)]^{i}\} = \sum_{x=0}^{L-1} (x - m_{1})^{i} p_{u}(x)$

Absolute central moments : $\hat{\mu}_{i} = E\{[\mu - E(u)]^{i}\} = \sum_{x=0}^{L-1} (x - m_{1})^{i} p_{u}(x)$

Entropy : $H = E[-\log_{2} p_{u}] = -\sum_{x=0}^{L-1} p_{u}(x) \log_{2} p_{u}(x)$ bits

4/1/2022 Special Topics in Image Proc.

12



Histogram Features[2]

- Mean
 - average value, general brightness
- Variance
 - tells us something about the contrast
- Skew
 - measures the asymmetry about the mean in the gray-level distribution
- Energy
 - tells us something about how the gray levels are distributed
 - maximum value : 1
 - the lager this value is, the easier it is to compress the image data



Histogram Features[3]

Entropy

a measure that tells us how many bits we need to code the image data

Second-order histogram features

- based on a joint probability distribution model
- provides statistics based on pairs of pixels and their corresponding gray levels
- be used in identifying texture within an image

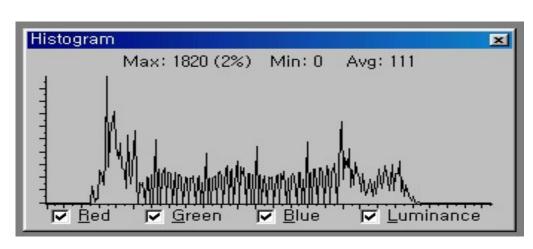
Histogram Features[4]

$$p_u(x_1, x_2) \equiv p_{u_1, u_2}(x_1, x_2) \equiv \text{Prob}[u_1 = x_1, u_2 = x_2]$$

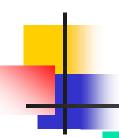
$$\approx \frac{\text{No. of pairs of pixels } u_1 = x_1, u_2 = x_2}{\text{Total no. of such pairs of pixels in the region}}$$







Special Topics in Image Proc.



Transform features[1]

Fourier Transform

$$F(u,v) = \frac{1}{N} \sum_{r=0}^{N-1} \sum_{c=0}^{N-1} I(r,c) e^{-j2\pi \frac{(ur+vc)}{N}}$$

$$MAGNITDE = |F(u,v)|$$

$$= \sqrt{[R(u,v)]^2 + [I(u,v)]^2}$$

$$PHASE = \Theta(u,v) = \tan^{-1} \left[\frac{I(u,v)}{R(u,v)} \right]$$

- magnitude : simply peak value
- phase: information about where objects are in an image,
 that is object's origin 4/1/2022 that is object's origin Pecial Topics in Image Proc.



Transform features[2]

Lowpass Filtering

- Pass low frequencies and attenuate or eliminates the high-frequency information.
- Used for <u>image compression</u> or for <u>hiding effects</u> <u>caused by noise.</u>
- Blur the image, although this blur is sometimes considered an enhancement.

Higpass Filtering

- Pass only high-frequency information.
- Can be used for edge enhancement.



Transform features[3]

- Bandpass & Bandreject Filters
 - Are specified by two cutoff frequencies → low and high cutoff
 - typically used in image restoration, enhancement and compression

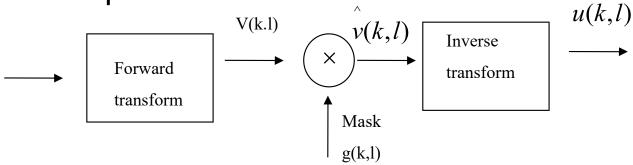
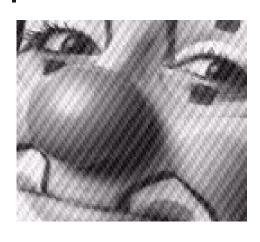
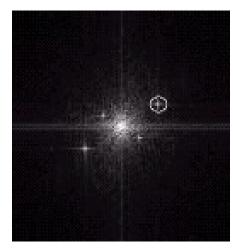


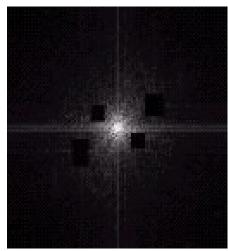
Fig.6 Transform feature extraction

Fourier Transform Example[4]









- (a) Before processing the image is marred by "noise".
- (b) An DFT of the image is displayed. Due to its regularity, the noise pattern stands out as four spikes. One of the spikes has been marked inside an AOI (Area of Interest) preparatory to its deletion.
- (c) The image is cleared.
- (d) The four noise spikes have been deleted from the DFT so that when the corrected DFT is applied.

4/1/2022

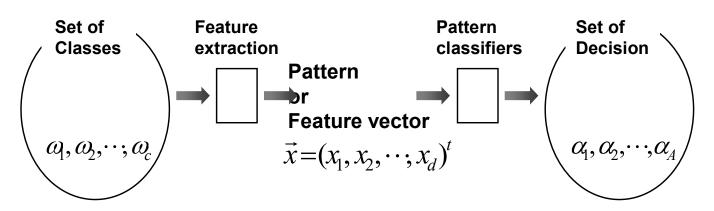
Special Topics in Image Proc.



2. Introduction to Classification

- Feature extraction

 classification
- Classification & segmentation
 - Closely related
 - Classification can lead to segmentation, and vice-versa



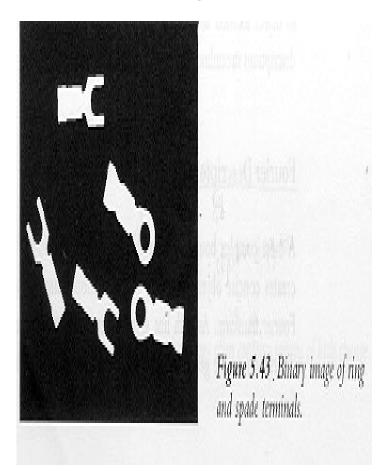
C: # of Class

A: # of Decisions

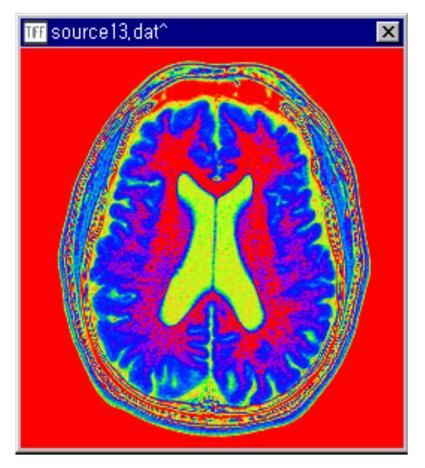
d: Dimension of feature vector

4

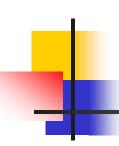
Hole in the Object?



MRI Classification



4/1/2022



Supervised vs. Unsupervised Classification

Supervised

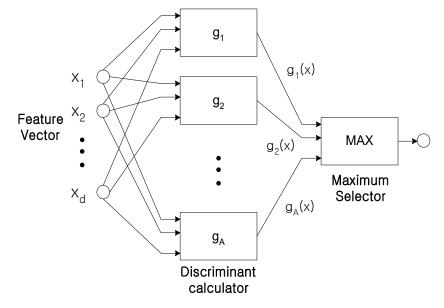
- Patterns from each class are available
- The system is "taught" to recognize patterns by means of various adaptive schemes
- The essentials of this approach
 - A set of training patterns of known classification
 - The implementation of an appropriate learning procedure

Unsupervised

- Only a set of training patterns of unknown classification are available
- Learning without a teacher

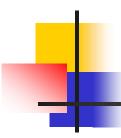
Distribution-free classification

- Do not require a priori probability distribution
 - Decision tree classification
 - Minimum distance classifier
 - K-nearest neighbor classifier
 - Linear discriminant function
- Decision rule
 - Discriminant function



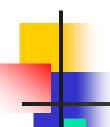
$$g_k(\mathbf{x}) > g_i(\mathbf{x})$$
 $k \neq i \Leftrightarrow \mathbf{x} \in S_k$

 $S_k = \text{Objects}(\text{patern classes})$



Statistical classification

- Bayes decision rule for minimum risk classifier
- Minimum error rate classifier
- Estimation of probability density function



Bayes' Decision Rule

- $P(\omega_i)$: a priori probability
- $P(\omega_i|\vec{x})$: a posteriori probability
- $P(\vec{x}|\omega_i)$: class conditional pdf

$$P(\omega_i | \vec{x}) = \frac{P(\omega_i, \vec{x})}{P(\vec{x})}$$
$$= \frac{P(\vec{x} | \omega_i) P(\omega_i)}{P(\vec{x})}$$

$$P(\vec{x}) = \sum_{i=1}^{c} P(\vec{x} | \omega_i) P(\omega_i)$$

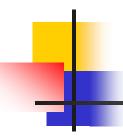


Bayes' decision rule(Two Class Problem)

2 class problem (i=1,2)

- a pixel belong to class 1 if p(x|1)p(1) > p(x|2)p(2)

- a pixel belong to class 2 if p(x|1) p(1)< p(x|2)p(2)



Unsupervised learning or clustering

Cluster

 A set of points in the feature space for which their <u>local</u> density is <u>large</u> compared to the density in the surrounding region

Clustering

- Useful for image segmentation and for classification of raw data to establish classes and prototypes
- Useful for vector quantization technique

General Clustering Algorithm

- Assume that we want to classify N samples X_1 , $X_2 \cdots X_N$
- Each sample is to be placed into one of L classes, $\omega_1, \omega_2, \cdots \omega_L$
- ullet The class to which the ith sample is assigned to $\omega_{ki} (i=1,\cdots N)$
- A classification is $oldsymbol{\Omega}$ a vector made up of the $oldsymbol{arOmega}_{ki}$ and

$$\Omega = \left[\omega_{ki} \cdots \omega_{kN}\right]^T$$

Clustering Criterion J is a function of Ω and X^* and can be written $J = J(\omega_{k1} \cdots \omega_{kN}; X_1 \cdots X_N) = J(\Omega; X^*)$

The best classification Ω satisfies max(J).

4/1/2022 Special Topics in Image Proc.



Similarity measure approach

 A general clustering algorithm is based on <u>split and</u> <u>merge</u> idea

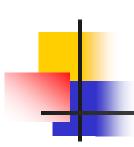
Basic

- The input vector are partitioned into subsets by using similarity measure
- Each partition is tested to check
- Subsets that are not sufficiently distinct are merged

Iterative(Isodata) method

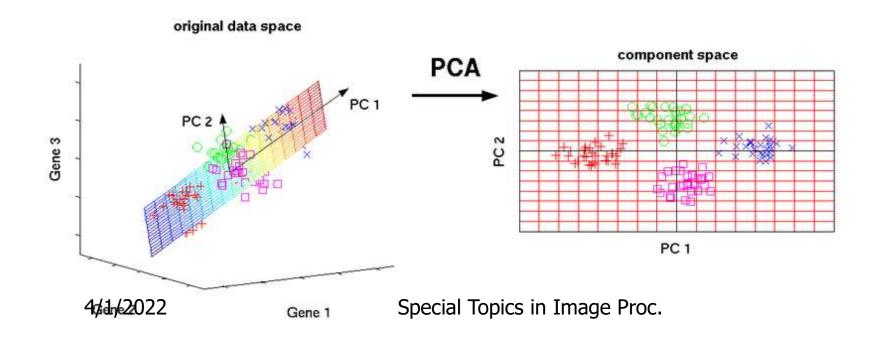
Basic ISODATA algorithm

- Choose some initial values for the means $\hat{\mu}_1, \hat{\mu}_2 \cdots \hat{\mu}_c$
- Classify the *n* samples by assigning them to the class of the closest mean
- Recompute the means as the average of the samples in their class.
- If any mean changed value, back to second procedure(reclassifying). Otherwise stop.

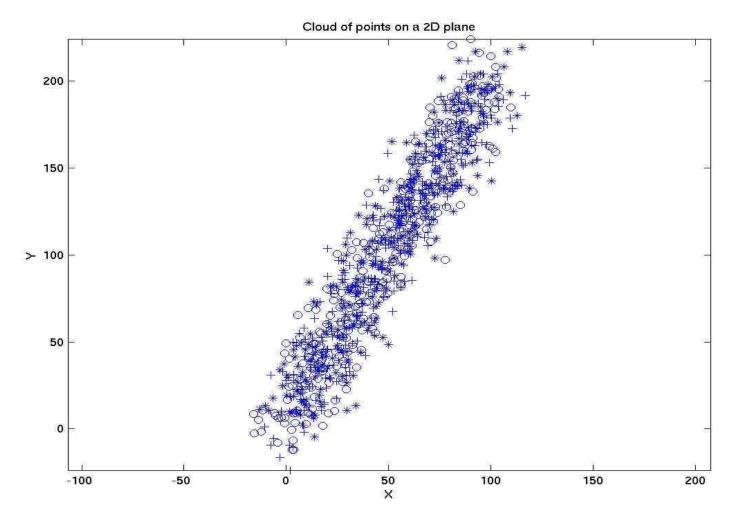


3. Principle Components Analysis (Karhunen-Loeve Transform)

• Principal component analysis (PCA) is a procedure for reducing the dimensionality of the variable space by representing it with a few orthogonal (uncorrelated) variables that capture most of its variability.



Suppose we are shown a set of points:





- We want to answer any of the following questions:
 - 1. What is the trend (direction) of the cloud of points
 - 2. Which is the direction of maximum variance (dispersion)
 - 3. Which is the direction of minimum variance (dispersion)
- Supposing I am only allowed to use a 1D description for this set of 2D points, i.e. using 1 coordinate instead of 2 coordinates (x,y). How should I represent the points in such a way that the overall error is minimized?
 - --- this is data compression problem

All these questions can be answered using Principle Components Analysis (PCA)

Principal Component Analysis (PCA)

Suppose we have N feature vectors

$$\{\mathbf{x}_1,\mathbf{x}_2,\mathbf{x}_3,\cdots,\mathbf{x}_N\}$$

We take the mean of the feature vectors

$$\overline{\mathbf{x}} = \frac{1}{N} \sum_{i=1}^{N} \mathbf{x}_{i}$$

Form the covariance matrix

$$C = \sum_{i=1}^{N} \mathbf{y}_{i} \mathbf{y}_{i}^{T}$$

where
$$\mathbf{y}_i = \mathbf{x}_i - \overline{\mathbf{x}}$$

Find the eigenvalues and eigenvectors of covariance matrix

$$C \mathbf{e}_i = \lambda_i \mathbf{e}_i$$

 λ_i where the **e**_i is a principal component with eigenvalue



- One problem with PCA when applied directly to the covariance matrix of the images is that the dimension of covariance matrix is huge!
- For example, for a 256x256 image, X will be of dimension 65536x1. So the covariance matrix of X will be of size 65536x65536 !!!
- To compute the eigenvector of such a huge matrix is clearly inconceivable.
- To solve this problem, Murakami and Kumar proposed a technique in the following paper:

Murakami, H. and Vijaya Kumar, B.V.K., "Efficient Calculation of Primary Images from a Set of Images", IEEE Trans. Pattern Analysis Special Topics in Image Proc. and Machine Intelligence, vol. PAMI-4, No. 5, Sep 1982.



To begin, we note that our goal is to find the eigenvectors and eigenvalues of the covariance matrix

$$C = \frac{1}{N} \sum_{i=1}^{N} (\mathbf{x}_i - \overline{\mathbf{x}}) (\mathbf{x}_i - \overline{\mathbf{x}})^T = \frac{1}{N} \sum_{i=1}^{N} \mathbf{y}_i \mathbf{y}_i^T$$

We will see later that instead of working directly on C, which is a huge matrix, we can find the eigenvectors and eigenvalues of an NxN matrix and then use it to deduce the eigenvectors and eigenvalues of $\,C\,$.

To see how this can be done, we first note that

$$\mathbf{y}_{i}\mathbf{y}_{i}^{T}\mathbf{e} = (\mathbf{y}_{i}^{T}\mathbf{e})\mathbf{y}_{i}$$

4

Therefore, if e is an eigenvector of C, then

$$\frac{1}{N} \sum_{i=1}^{N} \mathbf{y}_{i} \mathbf{y}_{i}^{T} \mathbf{e} = \lambda \mathbf{e}$$

$$\Rightarrow \frac{1}{N} \sum_{i=1}^{N} (\mathbf{y}_{i}^{T} \mathbf{e}) \mathbf{y}_{i} = \lambda \mathbf{e}$$

$$\Rightarrow \frac{1}{N} \sum_{i=1}^{N} a_{i} \mathbf{y}_{i} = \lambda \mathbf{e}$$
Let $b_{i} = \frac{a_{i}}{\lambda N}$
then $\mathbf{e} = \sum_{i=1}^{N} b_{i} \mathbf{y}_{i}$

So if $\mathbf e$ is an eigenvector of C, $\mathbf e$ can be expressed as a

Therefore,
$$\frac{1}{N} \sum_{i=1}^{N} \mathbf{y}_{i} \mathbf{y}_{i}^{T} \mathbf{e} = \lambda \sum_{i=1}^{N} b_{i} \mathbf{y}_{i}$$

$$\Rightarrow \frac{1}{N} \sum_{i=1}^{N} \mathbf{y}_{i} \mathbf{y}_{i}^{T} \left(\sum_{j=1}^{N} b_{j} \mathbf{y}_{j} \right) = \lambda \sum_{i=1}^{N} b_{i} \mathbf{y}_{i}$$

$$\Rightarrow \frac{1}{N} \sum_{i=1}^{N} \mathbf{y}_{i} \sum_{j=1}^{N} b_{j} \mathbf{y}_{i}^{T} \mathbf{y}_{j} = \lambda \sum_{i=1}^{N} b_{i} \mathbf{y}_{i}$$

$$\Rightarrow \frac{1}{N} \sum_{j=1}^{N} \mathbf{y}_{i}^{T} \mathbf{y}_{j} b_{j} = \lambda b_{i}$$

Therefore, the vector
$$\mathbf{b} = \begin{bmatrix} b_1 \\ b_2 \\ \vdots \\ b_N \end{bmatrix}$$
 is the eigenvector of the matrix

with (i,j) elements given by
$$\begin{cases} \frac{1}{4} y_{j}^{T} y_{j} \\ \text{Special Topics in Image Proc.} \end{cases}$$

So the eigenvector of the matrix $\frac{1}{N}\sum_{i=1}^{N}\mathbf{y}_{i}\mathbf{y}_{i}^{T}$, i.e. e, is given by

$$\mathbf{e} = \sum_{i=1}^{N} b_i \mathbf{y}_i$$

$$= [\mathbf{y}_1 \quad \mathbf{y}_2 \quad \cdots \quad \mathbf{y}_N] \mathbf{b}$$
The eigenvalues of the matrix with elements.

Note that the eigenvalues of the matrix with elements $\left\{\frac{1}{N}\mathbf{y}_{i}^{T}\mathbf{y}_{j}\right\}$ Note that the eigenvalues of the matrix $\frac{1}{N}\sum_{i=1}^{N}\mathbf{y}_{i}\mathbf{y}_{i}^{T}$

$$\left\{\frac{1}{N}\mathbf{y}_i^T\mathbf{y}_j\right\}$$

In conclusion, the task of finding the eigenvectors of a huge matrix of size $M \times M$ (where M is the number of pixels in an image), is reduced to the task of finding the eigenvectors of a NxN matrix.



Suppose we have N feature vectors $\{\mathbf{x}_1, \mathbf{x}_2, \mathbf{x}_3, \dots, \mathbf{x}_N\}$ We want to find the principle components. Since a direct computation of covariance matrix is not feasible due to huge dimension of covariance matrix, we do the following steps instead:

Step 1:

We take the mean of the feature vectors

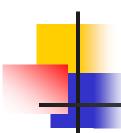
$$\overline{\mathbf{x}} = \frac{1}{N} \sum_{i=1}^{N} \mathbf{x}_{i}$$

Step 2:

Subtract the mean \bar{x} from each of the feature vectors

$$\mathbf{X}_{i}$$

$$\mathbf{y}_i = \mathbf{x}_i - \overline{\mathbf{x}}$$
Special Topics in Image Proc.



Summary(2)

Step 3:

Form the inner product matrix

$$M = \begin{bmatrix} \mathbf{y}_1^T \mathbf{y}_1 & \mathbf{y}_1^T \mathbf{y}_2 & \cdots & \mathbf{y}_1^T \mathbf{y}_N \\ \mathbf{y}_2^T \mathbf{y}_1 & \mathbf{y}_2^T \mathbf{y}_2 & \cdots & \mathbf{y}_2^T \mathbf{y}_N \\ \vdots & \vdots & \vdots & \vdots \\ \mathbf{y}_N^T \mathbf{y}_1 & \mathbf{y}_N^T \mathbf{y}_2 & \cdots & \mathbf{y}_N^T \mathbf{y}_N \end{bmatrix}$$
ep 4:

Step 4:

Find the eigenvectors and eigenvalues of M

$$\mathbf{M} \mathbf{b} = \mathbf{g}_{\mathsf{pecial}} \mathbf{b}_{\mathsf{Topics in Image Proc.}}$$

4

Summary(3)

Step 5:

If
$$[\mathbf{e}_1 \ \mathbf{e}_2 \ \mathbf{e}_3 \cdots \mathbf{e}_N]$$
 and $[\lambda_1 \ \lambda_2 \ \lambda_3 \cdots \lambda_N]$

are the principle components and the corresponding

eigenvalues respectively of $\begin{bmatrix} \mathbf{x}_1 \ \mathbf{x}_2 \ \mathbf{x}_3 \cdots \mathbf{x}_N \end{bmatrix}$

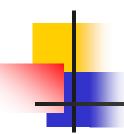
then

$$\mathbf{e}_i = \left[\mathbf{y}_1 \ \mathbf{y}_2 \ \mathbf{y}_3 \cdots \mathbf{y}_N \right] \mathbf{b}_i$$

$$\lambda_i = \alpha_i$$



Pattern Recognition (2)



Pattern Recognition

Pattern recognition is:

- 1. A research area in which patterns in data are found, recognized, discovered, ...whatever.
- 2. A catchall phrase that includes
 - classification
 - clustering
 - data mining
 -



Two Schools of Thought

1. Statistical Pattern Recognition

The data is reduced to vectors of numbers and statistical techniques are used for the tasks to be performed.

2. Structural Pattern Recognition

The data is converted to a discrete structure (such as a grammar or a graph) and the techniques are related to computer science subjects (such as parsing and graph matching).

1

In this class

- 1. How should objects to be classified be represented?
- 2. What algorithms can be used for recognition (or matching)?
- 3. How should learning (training) be done?



Classification in Statistical PR

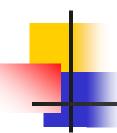
- A class is a set of objects having some important properties in common
- A feature extractor is a program that inputs the data (image) and extracts features that can be used in classification.
- A classifier is a program that inputs the feature vector and assigns it to one of a set of designated classes or to the "reject" class.

Feature Vector Representation

- X=[x1, x2, ..., xn], each xj a real number
- xj may be an object measurement
- xj may be count of object parts
- Example: object rep. [#holes, #strokes, moments, ...]

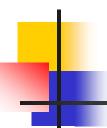
Possible features for char recog.

(class) character	area	height	vidth	number #holes	number #strokes	(cr,cy)	best axis	least inertia
' A '	medium	high	3/4	1	3	1/2,2/3	90	medium
'B'	medium	dgid	3/4	2	1	1/3,1/2	90	large
'8'	medium	dgid	2/3	2	0	1/2,1/2	90	medium
'O'	medium	dgid	2/3	1	0	1/2,1/2	90	large
,1,	lov	high	1/4	0	1	1/2,1/2	90	lop
' !	Agid	dgid	1	0	4	1/2,2/3	90	large
, I ,	dgid	high	3/4	0	2	1/2,1/2	7	large
2 * 2	medium	lop	1/2	0	0	1/2,1/2	?	large
7 _ 7	109	100	2/3	D	1	1/2,1/2	0	105
7/2	lov	high	2/3	0	1	1/2,1/2	60	109



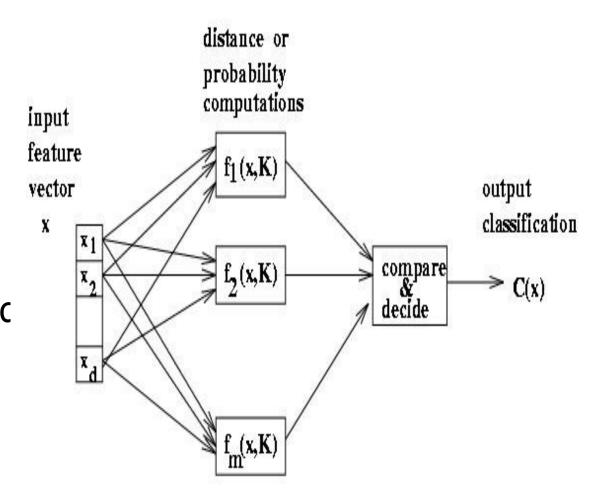
Some Terminology

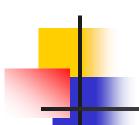
- Classes: set of m known categories of objects
 - (a) might have a known description for each
 - (b) might have a set of samples for each
- Reject Class:
 - a generic class for objects not in any of the designated known classes
- Classifier: Assigns object to a class based on features



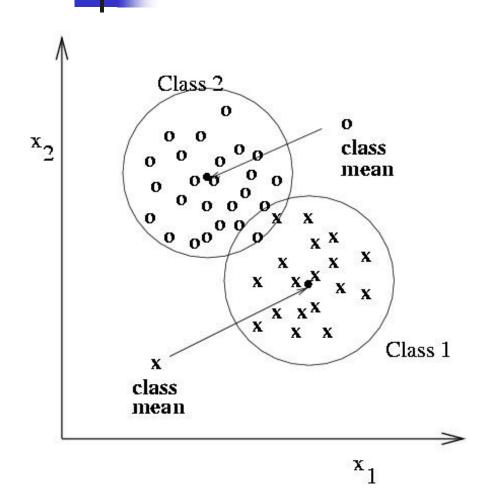
Discriminant functions

- Functions f(x, K)
 perform some
 computation on
 feature vector x
- Knowledge K from training or programming is used
- Final stage determines class





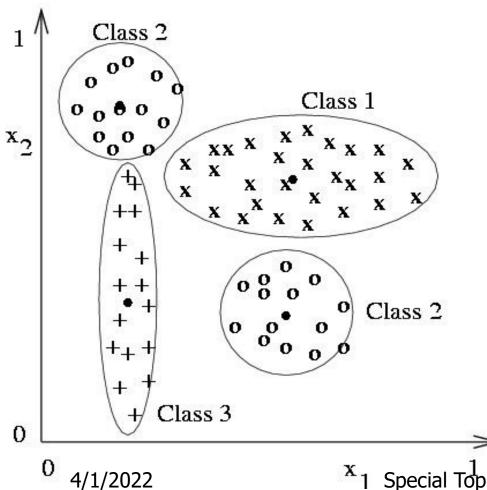
Classification using nearest class mean



- Compute the Euclidean distance between feature vector X and the mean of each class.
- Choose closest class, if close enough (reject otherwise)



Nearest mean might yield poor results with complex structure



Class 2 has two modes; where is its mean?

But if modes are detected, two subclass mean vectors can be used

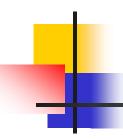


Scaling coordinates by std dev

We can compute a modified distance from feature vector \mathbf{x} to class mean vector \mathbf{x}_c by scaling by the spread, or standard deviation, σ_i of class c along each dimension i.

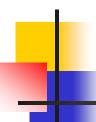
scaled Euclidean distance from x to class mean
$$x_c$$
: $\parallel x - x_c \parallel = \sqrt{\sum_{i=1,d}((\mathbf{x}[i] - \mathbf{x}_c[i])/\sigma_i)^2}$

In the previous 3 class problem, an observed X near the top of the Class 3 distribution will scale to be closer to the mean of Class 3 than to the mean of Class 2. Without scaling, X would be closer to the mean of Class 2.



Nearest Neighbor Classification

- Keep all the training samples in some efficient look-up structure.
- Find the nearest neighbor of the feature vector to be classified and assign the class of the neighbor.
- Can be extended to K nearest neighbors.



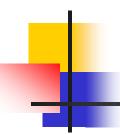
Bayesian decision-making

- classify into class ω_i that is most likely based on observations X
- In order to compute the likelihoods given the measurement X, the following distributions are needed.

class conditional distribution:
$$p(x|\omega_i)$$
 for each class omega(1)
a priori probability: $P(\omega_i)$ for each class ω_i (2)
unconditional distribution: $p(x)$ (3)

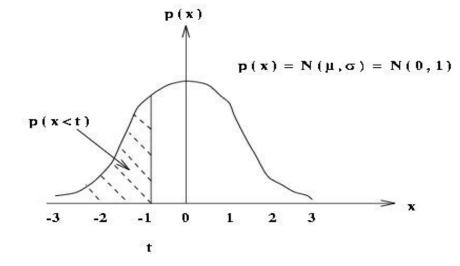
• use Bayes rule if all of the classes ω_i are disjoint

$$P(\omega_i|x) = \frac{p(x|\omega_i)P(\omega_i)}{p(x)} = \frac{p(x|\omega_i)P(\omega_i)}{\sum_{i=1,m} p(x|\omega_i)P(\omega_i)}$$
(4)



Normal distribution

- 0 mean and unit std deviation
- Table enables us to fit histograms and represent them simply
- New observation of variable x can then be translated into probability



t	p(x <t)< th=""><th>t</th><th>p(x<t)< th=""><th>t</th><th>p(x<t)< th=""></t)<></th></t)<></th></t)<>	t	p(x <t)< th=""><th>t</th><th>p(x<t)< th=""></t)<></th></t)<>	t	p(x <t)< th=""></t)<>
-3.0	0.0014	-2.0	0.0227	-1.0	0.1587
-2.9	0.0019	-1.9	0.0287	-0.9	0.1841
-2.8	0.0026	-1.8	0.0359	-0.8	0.2119
-2.7	0.0035	-1.7	0.0446	-0.7	0.2420
-2.6	0.0047	-1.6	0.0548	-0.6	0.2743
-2.5	0.0062	-1.5	0.0668	-0.5	0.3085
-2.4	0.0082	-1.4	0.0808	-0.4	0.3446
-2.3	0.0107	-1.3	0.0968	-0.3	0.3821
-2.2	0.0139	-1.2	0.1151	-0.2	0.4207
-2.1	0.0179	-1.1	0.1357	-0.1	0.4602
				0.0	0.5000

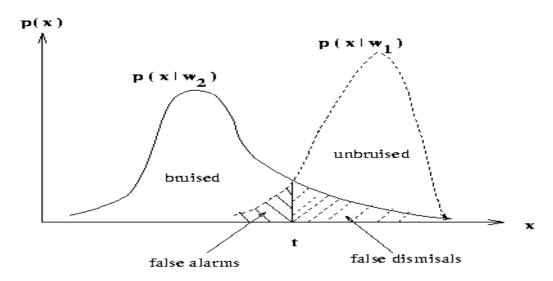


Parametric Models can be used

Parametric Models for Distributions

A normal distribution characterized by mean μ and standard deviation σ is defined as follows.

$$p(x) = N(\mu, \sigma)(x) = \frac{1}{\sqrt{2\pi}\sigma} exp[-\frac{1}{2}(\frac{x-\mu}{\sigma})^2]$$

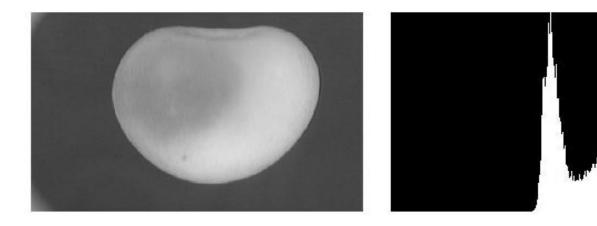


Distributions for intensity measurement x conditioned on whether x is taken from an unbruised or bruised cherry.

Cherry with bruise

- Intensities at about 750 nanometers wavelength
- Some overlap caused by cherry surface turning away

Cherry with bruise and histogram



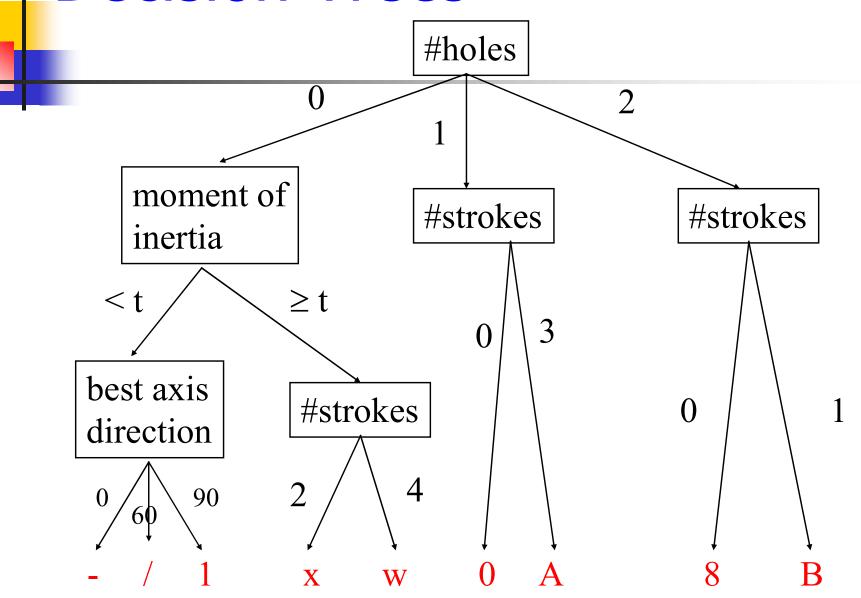
Cherry with bruise and histogram of intensities at about 750 nanometers wavelength. Some of the overlap is explained by unbruised surface that curves away from the viewpoint and thus appears dark like bruised tissue facing the viewpoint.



Classifiers often used in CV

- Decision Tree Classifiers
- Artificial Neural Net Classifiers
- Bayesian Classifiers and Bayesian Networks (Graphical Models)
- Support Vector Machines

Decision Trees





Decision Tree Characteristics

- 1. Training: How do you construct one from training data? Entropy-based Methods
- Strengths: Easy to Understand
- 3. Weaknesses: Overtraining

Entropy

Given a set of training vectors S, if there are c classes,

Entropy(S) =
$$\sum_{i=1}^{c}$$
 -pi log₂(pi)

Where pi is the proportion of category i examples in S.

If all examples belong to the same category, the entropy is 0.

If the examples are equally mixed (1/c examples of each class), the entropy is a maximum at 1.0.

e.g. for
$$c=2$$
, $-.5 \log_2 .5 - .5 \log_2 .5 = -.5(-1) -.5(-1) = 1$

4/1/2022 Special Topics in Image Proc.

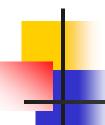
Information Gain

The information gain of an attribute A is the expected reduction in entropy caused by partitioning on this attribute.

Gain(S,A) = Entropy(S) -
$$\sum_{v \in Values(A)} \frac{|Sv|}{|S|}$$
 Entropy(Sv)

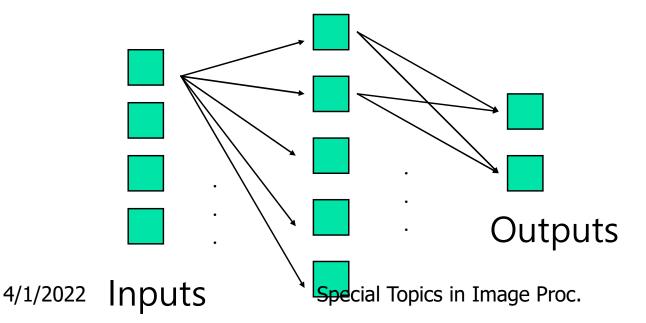
where Sv is the subset of S for which attribute A has value v.

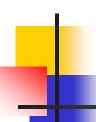
Choose the attribute A that gives the maximum information gain.



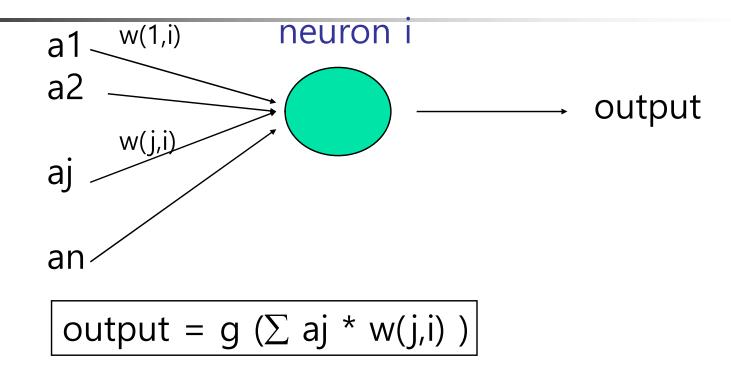
Artificial Neural Nets

- Artificial Neural Nets (ANNs) are networks of artificial neuron nodes, each of which computes a simple function.
- An ANN has an input layer, an output layer, and "hidden" layers of nodes.





Node Functions



Function g is commonly a step function, sign function, or sigmoid function (see text).



Neural Net Learning

That's beyond the scope of this text; only simple feed-forward learning is covered.

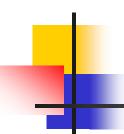
The most common method is called back propagation.

Support Vector Machines (SVM)

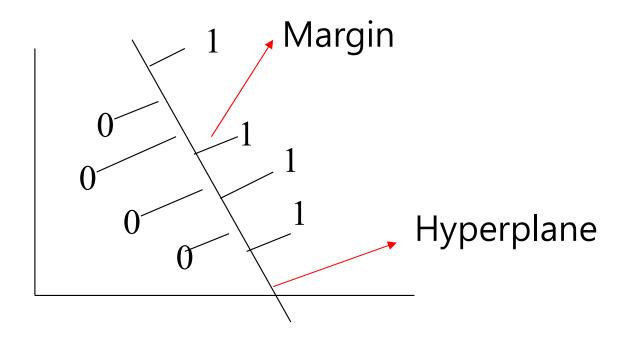
Support vector machines are learning algorithms that try to find a hyperplane that separates the differently classified data the most.

They are based on two key ideas:

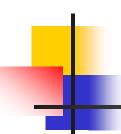
- Maximum margin hyperplanes
- A kernel 'trick'.



Maximal Margin



Find the hyperplane with maximal margin for all the points. This originates an optimization problem Which has a unique solution (convex problem).

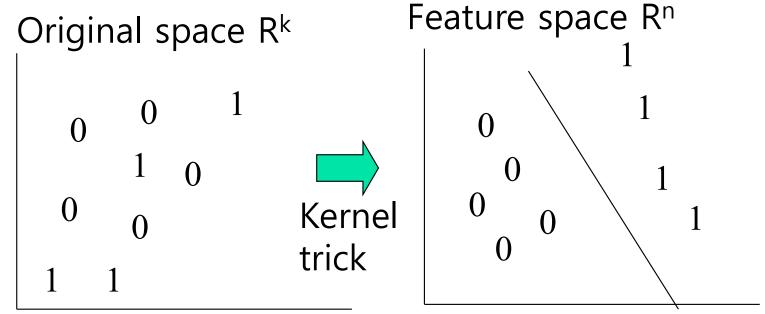


Non-separable data

What can be done if data cannot be separated with a hyperplane?

The kernel trick

The SVM algorithm implicitly maps the original data to a feature space of possibly infinite dimension in which data (which is not separable in the original space) becomes separable in the feature space.

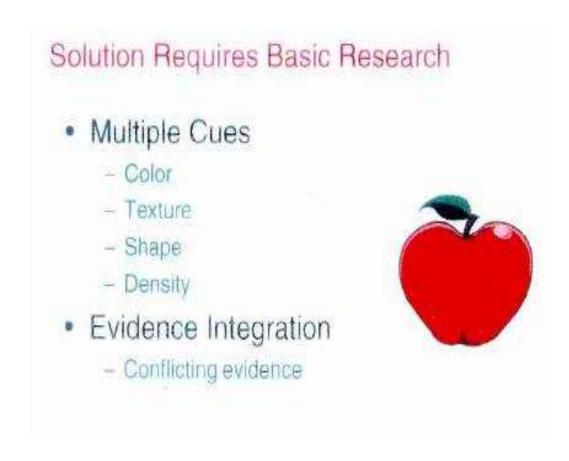


Veggie Vision by IBM

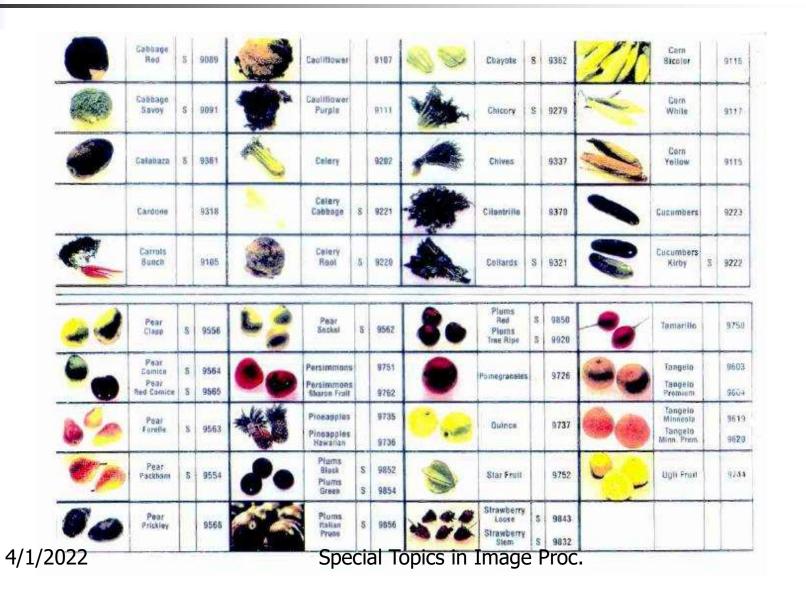
Ideas about a practical system to make more efficient the selling and inventory of produce in a grocery store.

10 years of R&D now

This information was shared by IBM researchers. Since that time, the system has been tested in small markets and has been modified according to that experience.



Up to 400 produce types



Practical problems of application environment

- Shrinkage -> Sweethearting
- Slow check-out
- Checker training
- Affixing PLUs
- Expensive packaging
- Solid waste

If asked, most grocers and supermarket owners would admit that they would rather not carry produce:



Rather than...



Engineering the solution

System to operate inside the usual checkout station

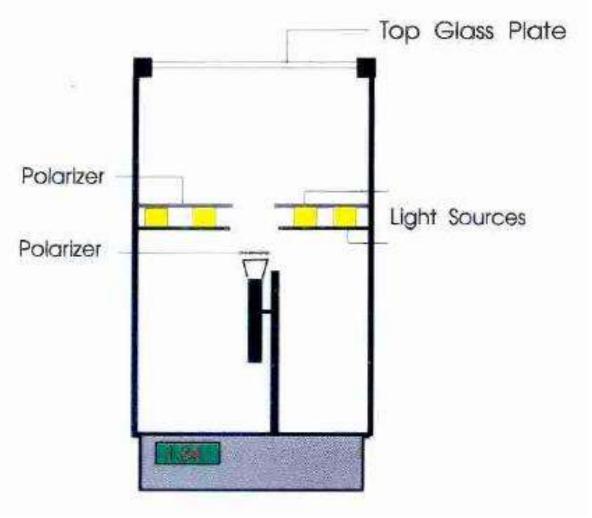
- together with bar code scanner
- together with scale
- together with accounting
- together with inventory
- together with employee
- within typical store environment



* figure shows system asking for help from the cashier in making final decision on touch screen



Modifying the scale

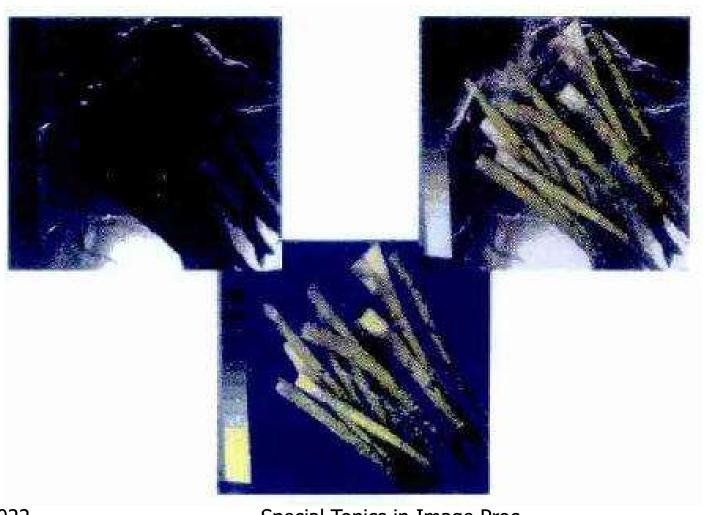


Need careful lighting engineering





Need to segment product from background





Quality segmented image obtained

Design of pattern recognition paradigm (from 1997)

FEATURES are: color, texture, shape, and size all represented uniformly by histograms

Matching

- signature = concatenated histograms
- multiple signatures per produce type
- · weighted absolute histogram difference
- · nearest neighbors retrieved
- "sure" if top 2 are the same type

Training

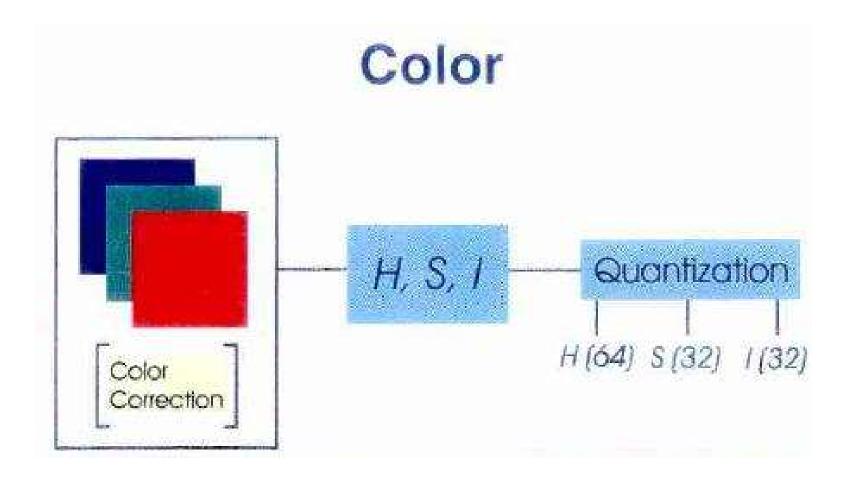
- · batch and incremental modes
- · store new signature of labelled example
- automatic selection skip if already "sure"
- signatures per class limited (e.g. 12)
- replace least useful and least recently used



Matching procedure

- Sample product represented by concatenated histograms: about 400 D
- 350 produce items x 10 samples = 3500 feature vectors of 400D each
- Have about 2 seconds to compare an unknown sample to 3500 stored samples (3500 dot products)
- Analyze the k nearest: if closest 2 are from one class, recognize that class (sure)

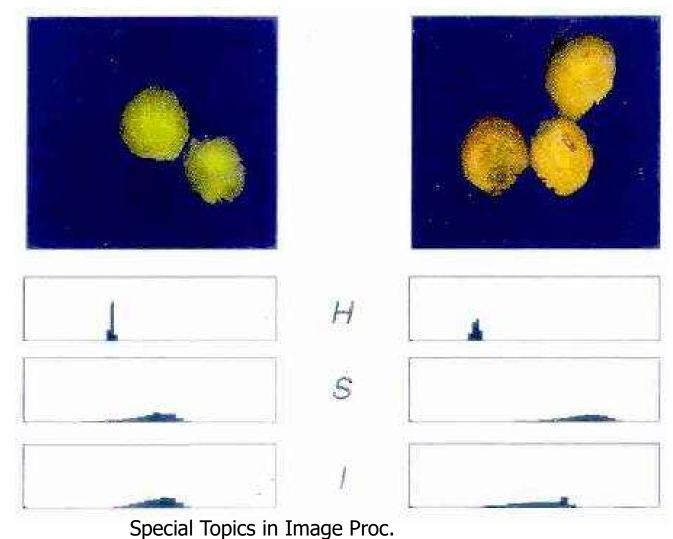
HSI for pixel color: 6 bits for hue, 5 for saturation and intensity



Histograms of 2 limes versus 3 lemons

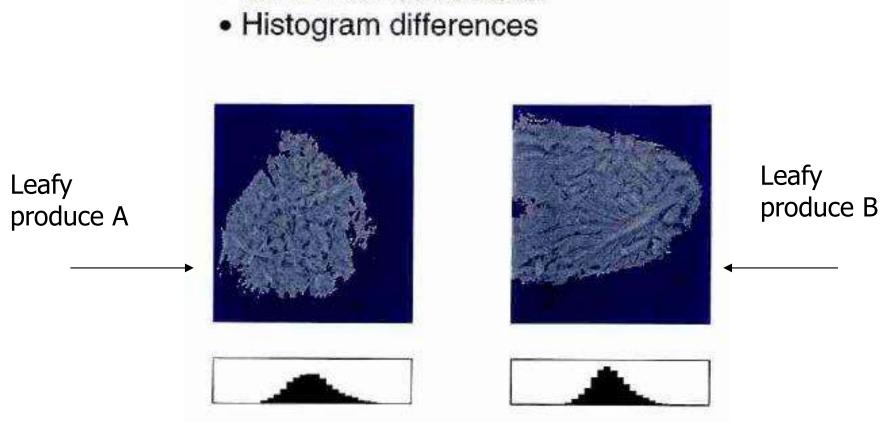
Distribution or population concept adds robustness:

- to size of objects
- to number of objects
- to small variations of color (texture, shape, size)



Texture: histogram results of LOG filter[s] on produce pixels

Center surround mask

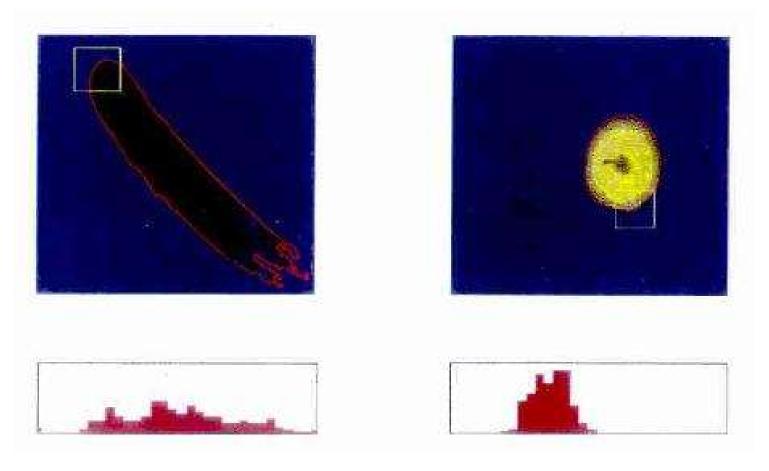


Shape: histogram of curvature of boundary of produce

Shape

- Chain-Coded Produce Boundary.
- Smooth Boundary.
- Compute Radii of Curvature Around Boundary.
- Histogram of Radii of Curvature.

Banana versus lemon or cucumber versus lime



Large range of curvatures

indicates complex object

4/1/2022

Special Topics in Image Proc.

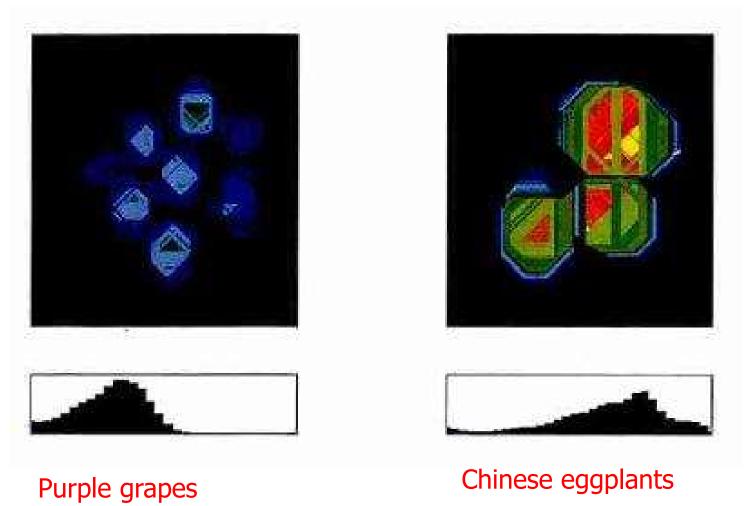
Small range of curvatures indicates roundish object



Size

- Threshold foreground image
- Find run lengths in 4 directions
- Label each pixel with minimum
- Histogram values and smooth

Each pixel gets a "size" as the minimum distance to boundary





Learning and adaptation

- System easy to train: show it produce samples and tell it the labels.
- During service: age out oldest sample; replace last used sample with newly identified one.
- When multiple labeled samples match the unknown, system asks cashier to select from the possible choices.