

# BE530 – Medical Deep Learning

– Representative GANs –

---

Byoung-Dai Lee

Division of AI Computer Science and Engineering

Kyonggi University

## ■ Perceptual Loss

- Perceptual Losses for Real-Time Style Transfer and Super Resolution (2016)

## ■ Multimodal Unsupervised Image-to-Image Translation (MUNIT)

- Multimodal Unsupervised Image-to-Image Translation (2018)

## ■ Cycle-Consistent Network

- CycleGAN: Unpaired Image-to-Image Translation using Cycle-Consistent Adversarial Networks (2017)

# Perceptual Loss

## ■ Key Objectives

- To train feed-forward transformation networks for image transformation tasks using **perceptual loss functions that depend on high-level features from a pre-trained loss network**

## ■ Perceptual loss function

- To measure high-level perceptual and semantic differences between images

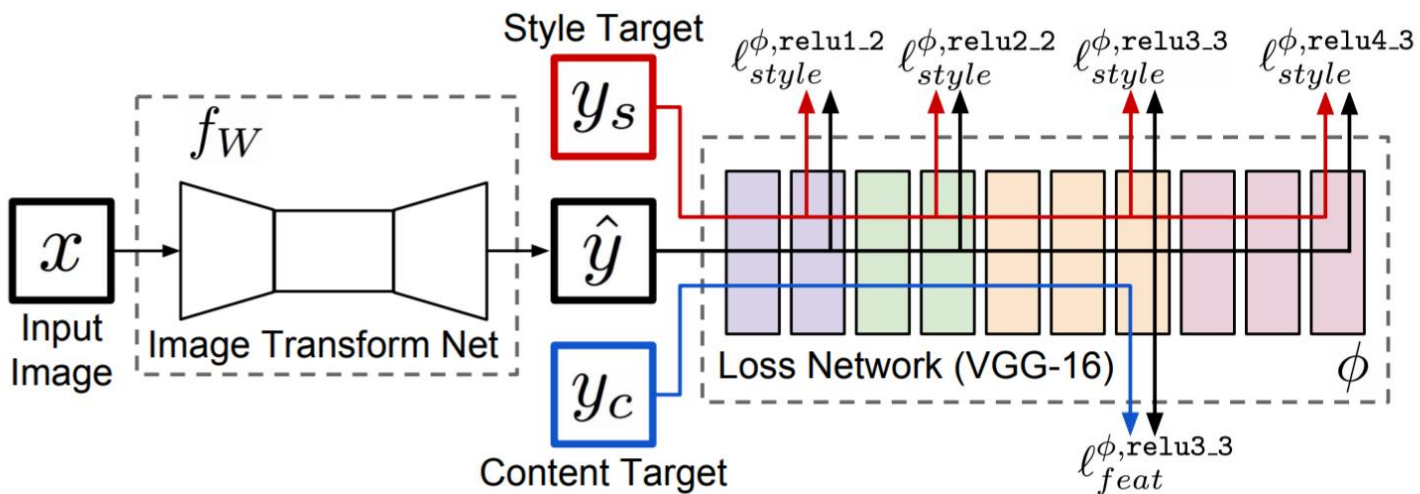
## ■ Insights

- Convolutional neural networks pretrained for image classification have already learned to encode the perceptual and semantic information we would like to measure in our loss functions

# Perceptual Loss (cont.)

## ■ System overview

- An image transformation network + a loss network



**Fig. 2.** System overview. We train an *image transformation network* to transform input images into output images. We use a *loss network* pretrained for image classification to define *perceptual loss functions* that measure perceptual differences in content and style between images. The loss network remains fixed during the training process.

# Perceptual Loss (cont.)

## ■ Loss network

The loss network  $\phi$  is used to define a *feature reconstruction loss*  $\ell_{feat}^\phi$  and a *style reconstruction loss*  $\ell_{style}^\phi$  that measure differences in content and style between images. For each input image  $x$  we have a *content target*  $y_c$  and a *style target*  $y_s$ . For style transfer, the content target  $y_c$  is the input image  $x$  and the output image  $\hat{y}$  should combine the content of  $x = y_c$  with the style of  $y_s$ ; we train one network per style target. For single-image super-resolution, the input image  $x$  is a low-resolution input, the content target  $y_c$  is the ground-truth high-resolution image, and the style reconstruction loss is not used; we train one network per super-resolution factor.

## ■ Image transformation network

- Style transfer
  - Input/Output images – 256×256×3 color images

# Perceptual Loss (cont.)

## ■ Perceptual loss function

### • Feature reconstruction loss

- To verify if the output image has similar feature representations
- The (squared, normalized) Euclidean distance between feature representations of the layer of the loss network

$$\ell_{feat}^{\phi,j}(\hat{y}, y) = \frac{1}{C_j H_j W_j} \|\phi_j(\hat{y}) - \phi_j(y)\|_2^2$$

$\phi_j(x)$  be the activations of the  $j$ th layer of the network

$$\hat{y} = f_W(x)$$

- **The feature reconstruction loss penalizes the output image when it deviates in content from the target image**
  - Using a feature reconstruction loss for training the image transformation networks encourages the output image to be perceptually similar to the target image
  - To minimize the feature reconstruction loss for early layers tends to produce images that are visually indistinguishable from the target image

# Perceptual Loss (cont.)

## ■ Perceptual loss function (cont.)

### • Style reconstruction loss

- It penalizes differences in style; colors, textures, common patterns, etc.

- Gram matrix

$G(i,j)$ 는 채널(i), 채널(j)의 feature map( $H \times W$ )간의 내적.  
따라서  $G(i,j)$ 가 큰 값을 가진다는 것은 두 채널이  
동시에 activation된다는 것을 의미함

$G_j^\phi(x)$  to be the  $C_j \times C_j$  matrix

$$G_j^\phi(x)_{c,c'} = \frac{1}{C_j H_j W_j} \sum_{h=1}^{H_j} \sum_{w=1}^{W_j} \phi_j(x)_{h,w,c} \phi_j(x)_{h,w,c'}$$

- Style reconstruction loss

$$\ell_{style}^{\phi,j}(\hat{y}, y) = \|G_j^\phi(\hat{y}) - G_j^\phi(y)\|_F^2$$

- Minimizing the style reconstruction loss preserves stylistic feature from the target image, but does not preserve its spatial structure

# Perceptual Loss (cont.)

## ■ Simple loss functions that depend only on low-level pixel information

- Pixel Loss

- The (normalized) Euclidean distance between the output image and the target image

$$\ell_{pixel}(\hat{y}, y) = \|\hat{y} - y\|_2^2 / CHW$$

- Total Variation Regularization

- To encourage spatial smoothness in the output image



## ■ Key Objective


- Unsupervised image-to-image translation
- Multimodal mapping
  - The model can generate diverse outputs from a given source domain image

## ■ Learning disentangled representation

- Disentangling content from style
  - Content – The underlying spatial structure
  - Style – The rendering of the structure

## ■ Assumptions

- The latent space of images can be decomposed into a content space and a style space
- Images in different domains share a common content space but not the style space
- The content code encodes the information that should be preserved during translation, while the style code represents remaining variations that are not contained in the input image



**Fig. 1.** An illustration of our method. (a) Images in each domain  $\mathcal{X}_i$  are encoded to a shared content space  $\mathcal{C}$  and a domain-specific style space  $\mathcal{S}_i$ . Each encoder has an inverse decoder omitted from this figure. (b) To translate an image in  $\mathcal{X}_1$  (e.g., a leopard) to  $\mathcal{X}_2$  (e.g., domestic cats), we recombine the content code of the input with a random style code in the target style space. Different style codes lead to different outputs.

# MUNIT (cont.)

## ■ Model

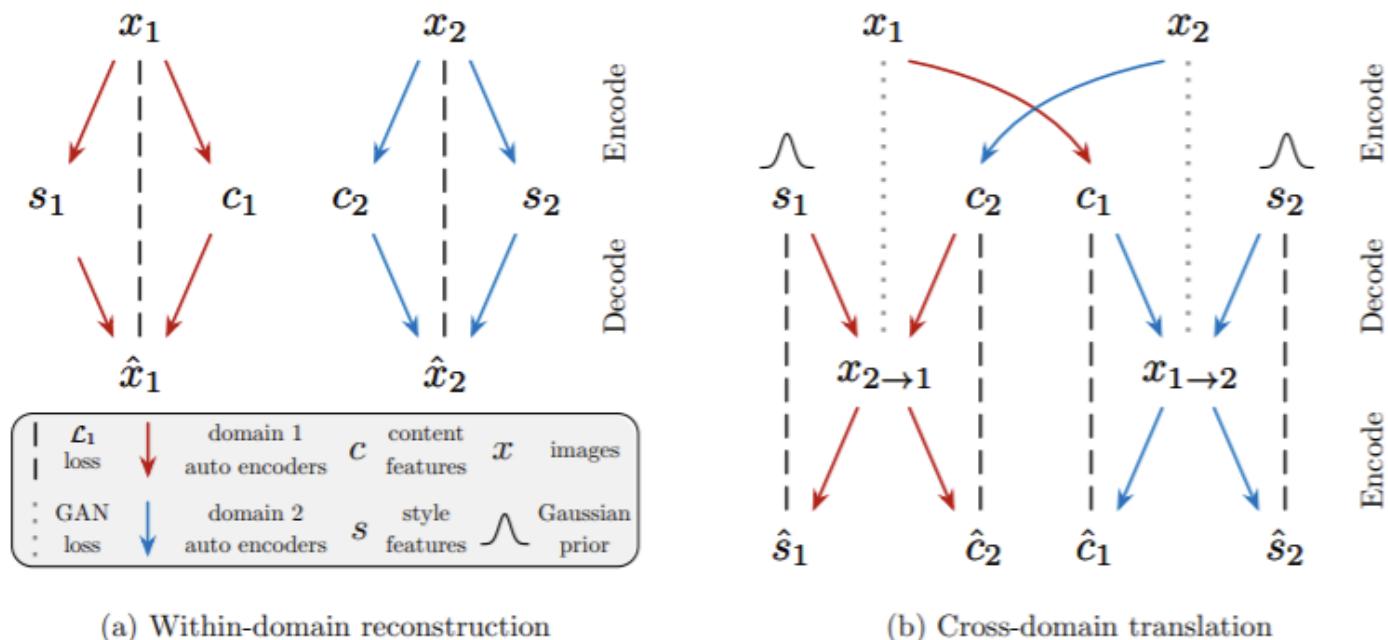


Fig. 2 shows an overview of our model and its learning process. Similar to Liu *et al.* [15], our translation model consists of an encoder  $E_i$  and a decoder  $G_i$  for each domain  $\mathcal{X}_i$  ( $i = 1, 2$ ). As shown in Fig. 2 (a), the latent code of each auto-encoder is factorized into a content code  $c_i$  and a style code  $s_i$ , where  $(c_i, s_i) = (E_i^c(x_i), E_i^s(x_i)) = E_i(x_i)$ . Image-to-image translation is performed by swapping encoder-decoder pairs, as illustrated in Fig. 2 (b). For example, to translate an image  $x_1 \in \mathcal{X}_1$  to  $\mathcal{X}_2$ , we first extract its content latent code  $c_1 = E_1^c(x_1)$  and randomly draw a style latent code  $s_2$  from the prior distribution  $q(s_2) \sim \mathcal{N}(0, \mathbf{I})$ . We then use  $G_2$  to produce the final output image  $x_{1 \rightarrow 2} = G_2(c_1, s_2)$ .

## ■ Loss function

- Bidirectional reconstruction loss

- Image reconstruction

$$\mathcal{L}_{\text{recon}}^{x_1} = \mathbb{E}_{x_1 \sim p(x_1)} [\|G_1(E_1^c(x_1), E_1^s(x_1)) - x_1\|_1]$$

- Latent reconstruction

$$\mathcal{L}_{\text{recon}}^{c_1} = \mathbb{E}_{c_1 \sim p(c_1), s_2 \sim q(s_2)} [\|E_2^c(G_2(c_1, s_2)) - c_1\|_1]$$

$$\mathcal{L}_{\text{recon}}^{s_2} = \mathbb{E}_{c_1 \sim p(c_1), s_2 \sim q(s_2)} [\|E_2^s(G_2(c_1, s_2)) - s_2\|_1]$$

- Adversarial loss

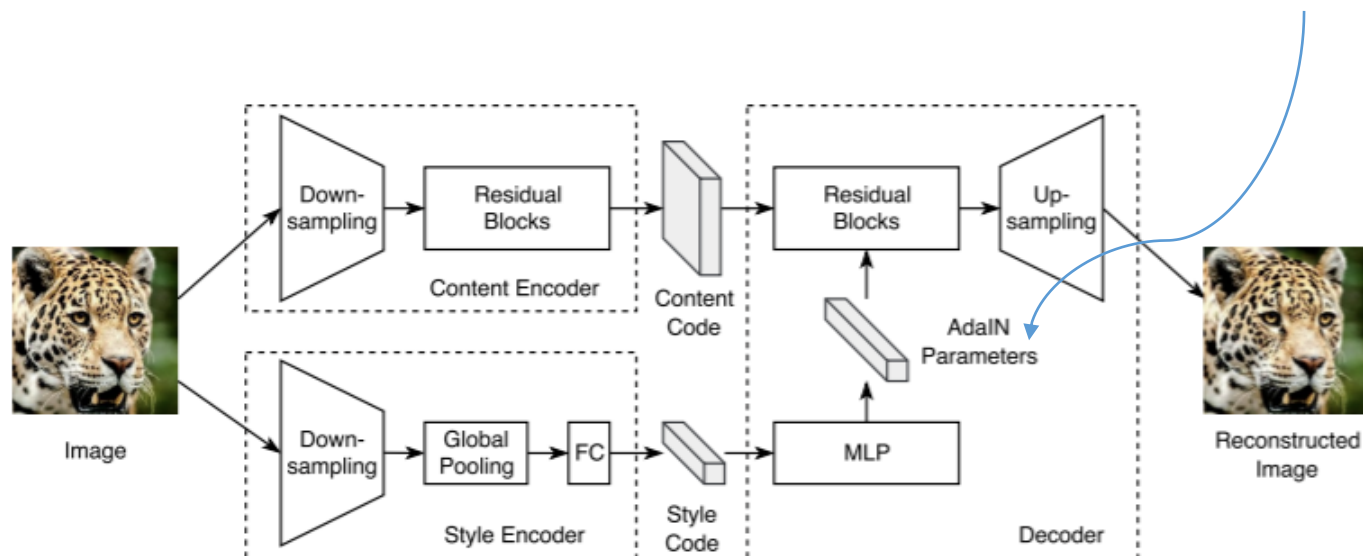
$$\mathcal{L}_{\text{GAN}}^{x_2} = \mathbb{E}_{c_1 \sim p(c_1), s_2 \sim q(s_2)} [\log(1 - D_2(G_2(c_1, s_2)))] + \mathbb{E}_{x_2 \sim p(x_2)} [\log D_2(x_2)]$$

- Total loss

$$\min_{E_1, E_2, G_1, G_2} \max_{D_1, D_2} \mathcal{L}(E_1, E_2, G_1, G_2, D_1, D_2) = \mathcal{L}_{\text{GAN}}^{x_1} + \mathcal{L}_{\text{GAN}}^{x_2} + \lambda_x(\mathcal{L}_{\text{recon}}^{x_1} + \mathcal{L}_{\text{recon}}^{x_2}) + \lambda_c(\mathcal{L}_{\text{recon}}^{c_1} + \mathcal{L}_{\text{recon}}^{c_2}) + \lambda_s(\mathcal{L}_{\text{recon}}^{s_1} + \mathcal{L}_{\text{recon}}^{s_2})$$

## ■ The architecture of auto-encoder

$$\text{AdaIN}(x, y) = \sigma(y) \left( \frac{x - \mu(x)}{\sigma(x)} \right) + \mu(y)$$



**Fig. 3.** Our auto-encoder architecture. The content encoder consists of several strided convolutional layers followed by residual blocks. The style encoder contains several strided convolutional layers followed by a global average pooling layer and a fully connected layer. The decoder uses a MLP to produce a set of AdaIN [54] parameters from the style code. The content code is then processed by residual blocks with AdaIN layers, and finally decoded to the image space by upsampling and convolutional layers.



# MUNIT (cont.)



(a) edges  $\leftrightarrow$  shoes



(b) edges  $\leftrightarrow$  handbags


# MUNIT (cont.)




(a) house cats  $\rightarrow$  big cats



(b) big cats  $\rightarrow$  house cats



(c) house cats  $\rightarrow$  dogs



(d) dogs  $\rightarrow$  house cats



(e) big cats  $\rightarrow$  dogs




(f) dogs  $\rightarrow$  big cats

# Cycle-Consistent GAN

## ■ Key Objective


- Capturing special characteristics of one image collection and figuring out how these characteristics could be translated into the other image collection, all in the absence of any paired training examples





# Cycle-Consistent GAN (cont.)

## ■ Model



## ■ Loss function

$$\begin{aligned} \mathcal{L}(G, F, D_X, D_Y) = & \mathcal{L}_{\text{GAN}}(G, D_Y, X, Y) \\ & + \mathcal{L}_{\text{GAN}}(F, D_X, Y, X) \\ & + \lambda \mathcal{L}_{\text{cyc}}(G, F), \end{aligned}$$

$$\begin{aligned} \mathcal{L}_{\text{GAN}}(G, D_Y, X, Y) = & \mathbb{E}_{y \sim p_{\text{data}}(y)} [\log D_Y(y)] \\ & + \mathbb{E}_{x \sim p_{\text{data}}(x)} [\log(1 - D_Y(G(x)))], \end{aligned}$$

$$\begin{aligned} \mathcal{L}_{\text{cyc}}(G, F) = & \mathbb{E}_{x \sim p_{\text{data}}(x)} [\|F(G(x)) - x\|_1] \\ & + \mathbb{E}_{y \sim p_{\text{data}}(y)} [\|G(F(y)) - y\|_1]. \end{aligned}$$

# Cycle-Consistent GAN (cont.)

## ■ Implementation

- Network architecture
  - Generator – based on the DNN of the perceptual losses paper
  - Discriminator – 70×70 PatchGAN
- Input image resolutions
  - 128×128 or 256×256

## ■ Training details

- To replace the negative log likelihood objective by a least-square loss

$$\mathcal{L}_{\text{GAN}}(G, D_Y, X, Y) = \mathbb{E}_{y \sim p_{\text{data}}(y)} [\log D_Y(y)] + \mathbb{E}_{x \sim p_{\text{data}}(x)} [\log(1 - D_Y(G(x)))]$$

→

we train the  $G$  to minimize  $\mathbb{E}_{x \sim p_{\text{data}}(x)} [(D(G(x)) - 1)^2]$   
and train the  $D$  to minimize  $\mathbb{E}_{y \sim p_{\text{data}}(y)} [(D(y) - 1)^2] + \mathbb{E}_{x \sim p_{\text{data}}(x)} [D(G(x))^2]$ .

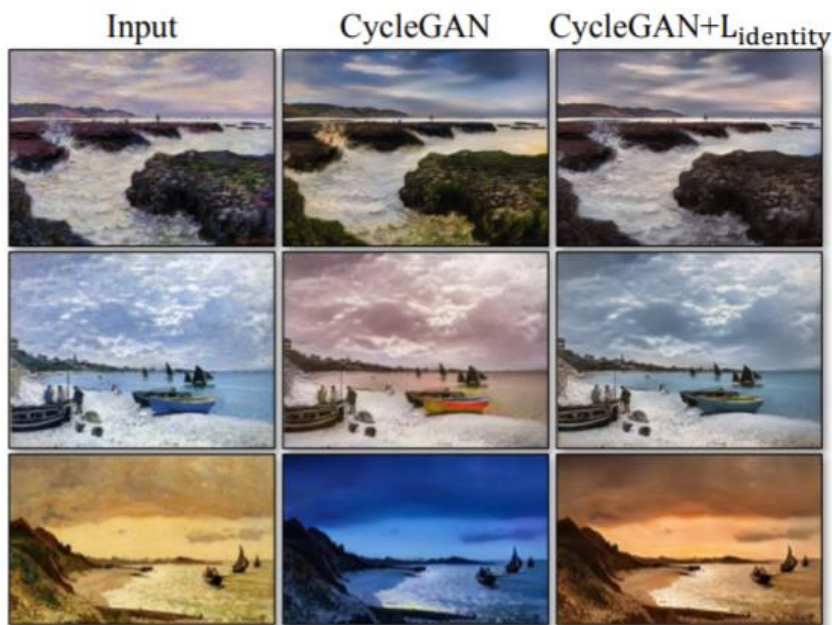
- To update the discriminator
  - Using a history of generated images rather than the ones produced by the latest generators
  - Previously created 50 images are kept in the buffer

# Cycle-Consistent GAN (cont.)

## ■ Identity Mapping

- $G(x) \rightarrow y$ 로 변환하는 맵핑의 경우,  $y$ 가  $G$ 의 입력으로 사용될 경우 새로운 값으로 변환하지 않고 원래 자신, 즉  $y$ 를 반환하게 함으로써 입력 이미지의 컬러값을 유지하도록 함

$$\mathcal{L}_{\text{identity}}(G, F) = \mathbb{E}_{y \sim p_{\text{data}}(y)} [\|G(y) - y\|_1] + \mathbb{E}_{x \sim p_{\text{data}}(x)} [\|F(x) - x\|_1]$$



# Cycle-Consistent GAN (cont.)

## ■ Limitations

- Color and texture changes → showed good performance
- Geometric changes → little success
- The distribution characteristics of the training datasets
  - “A person riding a horse” 이미지가 학습데이터에 없을 경우 유사한 이미지가 입력으로 제공될 경우 translation되는 이미지의 품질은 나쁠 수 있음

# References

- Perceptual Losses for Real-Time Style Transfer and Super Resolution, <https://arxiv.org/abs/1603.08155>
- CycleGAN: Unpaired Image-to-Image Translation using Cycle-Consistent Adversarial Networks, <https://arxiv.org/abs/1703.10593>
- Multimodal Unsupervised Image-to-Image Translation, <https://arxiv.org/pdf/1804.04732.pdf>
- Arbitrary Style Transfer in Real-Time With Adaptive Instance Normalization, [https://openaccess.thecvf.com/content\\_ICCV\\_2017/papers/Huang\\_Arbitrary\\_Style\\_Transfer\\_ICCV\\_2017\\_paper.pdf](https://openaccess.thecvf.com/content_ICCV_2017/papers/Huang_Arbitrary_Style_Transfer_ICCV_2017_paper.pdf)

**ANY  
QUESTIONS?**